

Proposal: Unsupervised Machine Translation with Word2Vec

II2202, Fall 2015

Bram Leenders
KTH, Royal Institute of Technology
Brinellvägen 8, 114 28 Stockholm
Sweden
b.c.leenders@gmail.com

Marc Romeyn
KTH, Royal Institute of Technology
Brinellvägen 8, 114 28 Stockholm
Sweden
marc.romeyn@gmail.com

ABSTRACT

Add an abstract or comment this part out

Keywords

Word2Vec, Machine Translation

1. INTRODUCTION

1.1 Problem Statement

The project will investigate how to...

1.2 Prior Work

Since the introduction of word2vec [3, 5] in 2013, the algorithm has seen a wide variety of usecases. In the initial paper [3], Mikolov et. al describe interesting relations between vectors corresponding to words. A famous example of how word2vec models relations between words as mathematical equations is $king - man + woman = queen$. The semantic relationships between man/woman and king/queen are preserved in the transformation of words to vectors, and can be expressed with basic algebra.

Subsequent papers have improved the algorithm both in terms of accuracy ([2]), performance, parallelization and extended the initial scope of applications. A good example of the latter is a paper by Boycheva [1], which uses word2vec not inside the natural language processing domain but to generate playlists. Based on a set of playlists, their word2vec-based algorithm can suggest new playlists with artists that go well together.

One of the applications of word2vec, is in exploiting similarities in languages for assistance in machine translation [6]. Mikolov et al. [4] released a subsequent paper on word2vec, in which they describe similarities between models for different languages. An example they give, is how the usage of the numbers one to five in English is very similar to how they are used in Spanish, and likewise for the names of animals. Figure 1 shows a graphical representation of word vectors in English and Spanish.

The similarities between languages can be used to predict translations for words without any human interaction or labeled input data. Using only unsupervised machine learning techniques, a computer could learn how to translate English to Spanish or vice versa. The only requirement is a large

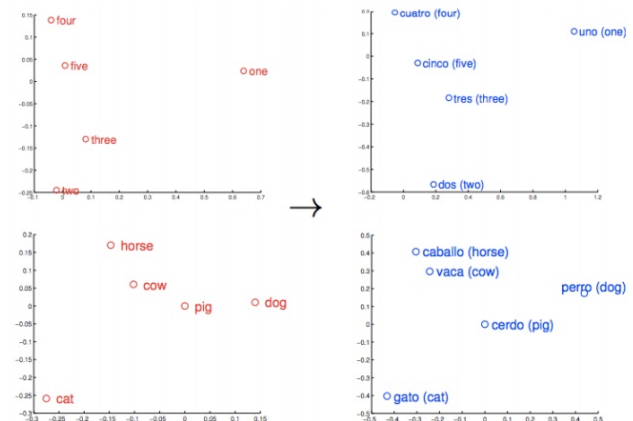


Figure 1: Vector representations of English and Spanish words, after dimensionality reduction and rotation. Notice the high level of similarity between both languages. Reprinted from Mikolov et al. [4]

amount of text in both languages to train the word2vec models on.

In this research, we will focus on this specific application of word2vec: using similarities in languages to provide translations of words.

It is important to note that word2vec only uses information of cooccurrences to model words. It does not learn grammatical concepts other than by statistical analysis. This limits our translation to single words; although the translator might be able to translate each word individually, it cannot learn that each finite verb must have a subject, that "we" is plural, etc. It will learn that "swim" is to "swimming" as "walk" is to "walking", but will not know that "swimming" is a gerund.

2. RESEARCH GOAL

Nunc sed pede. Praesent vitae lectus. Praesent neque justo, vehicula eget, interdum id, facilisis et, nibh. Phasellus at purus et libero lacinia dictum. Fusce aliquet. Nulla eu ante placerat leo semper dictum. Mauris metus. Curabitur lobortis. Curabitur sollicitudin hendrerit nunc. Donec ultrices lacus id ipsum.

2.1 Research Purpose

Abstract description of goals

Nunc sed pede. Praesent vitae lectus. Praesent neque justo, vehicula eget, interdum id, facilisis et, nibh. Phasellus at purus et libero lacinia dictum. Fusce aliquet. Nulla eu ante placerat leo semper dictum. Mauris metus. Curabitur lobortis. Curabitur sollicitudin hendrerit nunc. Donec ultrices lacus id ipsum.

2.2 Goals

Pretty much what our deliverables are...

2.3 Tasks

What do we have to do to achieve this?

- Build (at least) two working setups for automated translation
 - For the single-model setup:
 - * Merge datasets, train single model
 - For the multi-model setup:
 - * Train individual models
 - * Train transformation matrix between models
- Make a verification mechanism
- Verify which setup works better

3. PLANNING

A global overview of the milestones we defined in our research is as follows;

Half September - October: Literary Study)

During this phase, a list of relevant papers (e.g. [2, 4, 6]) is collected, short-listed to a readable size and read. Each paper will get a brief (informal) summarization to capture the essence of the paper insofar as it is relevant to this research. These summaries help process the information and provide quick access during the practical research.

Not all papers relevant to the research will be read during this phase, so next phases will include a fair share of reading.

October: Preparing experiments

In this phase, we gather the required corpora (see Section 5) and implement our translation programs.

November: Running experiments

Running the experiments includes training the models (which might take several days) and running the translation programs against a set of predefined correct translations.

Before 12 December: Incorporating results in paper

During this phase, we will incorporate the results obtained in the experiments into our final paper.

4. ALLOCATION OF RESPONSIBILITIES

Marc pledges to provide sufficient beer. Bram will take care of coffee and possibly some Club-Maté.

WRITE SOME ACTUAL STUFF HERE

5. RESOURCES

We have identified three main resources needed for this research: a large English corpus, a large Dutch corpus and computational resources to analyze the data.

The corpora for both languages should be several tens of gigabytes or more. Mikolov et al. [5] specifically state that a large amount of training data is crucial for word2vec to build a correct model. They speak of corpora 30 billion words, with a significant decrease in quality of the results when lowering the corpus size to only 6 billion words.

For this research, we will look at the following corpora:

- *Reddit corpus*¹ (1TB uncompressed JSON, 50 billion words) Corpus containing all posts on Reddit. It is mostly English, but also contains a few other languages.
- *Wikipedia EN crawl*² (60GB uncompressed XML, 3 billion words) Corpus of all text on the English Wikipedia (ignoring revisions).
- *Wikipedia NL Crawl*³ (5GB uncompressed XML, 250 million words) Corpus of all text on the Dutch Wikipedia (ignoring revisions).
- *SoNaR-500*⁴ (60GB uncompressed text, >500 million words) Curated corpus gathered by several Dutch universities, consisting of newspapers, subtitles, etc.

These datasets are all open for academic usage, in various degrees of openness. The Reddit corpus is available over bit-torrent, the Wikipedia data is freely available for download under CC-BY-SA license, and SoNaR-500 is available after signing a special research license agreement.

Before processing the corpora with word2vec, we will convert them to a universal format: plain text, with special characters (everything that is not alphanumeric, space or dash) removed, and one entity (post, article or article) per line.

To make our research reproducible, we will publish all our code online and provide lists with hashes of all our input and intermediate steps (e.g. after cleaning the data).

6. REFERENCES

- [1] N. Boycheva. Distributional similarity music recommendations versus spotify: A comparison based on user evaluation. 2015.

¹https://archive.org/details/2015_reddit_comments_corpus

²<https://dumps.wikimedia.org/enwiki/20150901/>

³<https://dumps.wikimedia.org/nlwiki/20150901/>

⁴<http://tst-centrale.org/producten/corpora/sonar-corpus/6-85>

- [2] O. Levy, Y. Goldberg, and I. Ramat-Gan. Linguistic regularities in sparse and explicit word representations. *CoNLL-2014*, page 171, 2014.
- [3] T. Mikolov, K. Chen, G. Corrado, and J. Dean. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*, 2013.
- [4] T. Mikolov, Q. V. Le, and I. Sutskever. Exploiting similarities among languages for machine translation. *arXiv preprint arXiv:1309.4168*, 2013.
- [5] T. Mikolov, I. Sutskever, K. Chen, G. S. Corrado, and J. Dean. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pages 3111–3119, 2013.
- [6] L. Wolf, Y. Hanani, K. Bar, and N. Dershowitz. Joint word2vec networks for bilingual semantic representations. *International Journal of Computational Linguistics and Applications*, 5(1):27–44, 2014.